

8fcd5a04-40fb-4260-89f9-ef0366c38900 (1)

June 10, 2021

1

2

2.1 CountVectorizer Edition

3

4 -

1 « »

- « » . , - .
.
.
.
F1 0.75.

1. .
2. .
3. .

BERT , .

`toxic_comments.csv`. *text* , *toxic* — .

1.1

```
[1]: !pip install catboost
```

Requirement already satisfied: catboost in c:\users\said\anaconda3\lib\site-packages (0.25.1)

Requirement already satisfied: pandas>=0.24.0 in c:\users\said\anaconda3\lib\site-packages (from catboost) (1.0.5)

Requirement already satisfied: six in c:\users\said\anaconda3\lib\site-packages (from catboost) (1.15.0)

Requirement already satisfied: graphviz in c:\users\said\anaconda3\lib\site-packages (from catboost) (0.16)

Requirement already satisfied: numpy>=1.16.0 in c:\users\said\anaconda3\lib\site-packages (from catboost) (1.18.5)

Requirement already satisfied: scipy in c:\users\said\anaconda3\lib\site-packages (from catboost) (1.5.0)

Requirement already satisfied: plotly in c:\users\said\anaconda3\lib\site-packages (from catboost) (4.9.0)

Requirement already satisfied: matplotlib in c:\users\said\anaconda3\lib\site-packages (from catboost) (3.2.2)

Requirement already satisfied: pytz>=2017.2 in c:\users\said\anaconda3\lib\site-packages (from pandas>=0.24.0->catboost) (2020.1)

Requirement already satisfied: python-dateutil>=2.6.1 in c:\users\said\anaconda3\lib\site-packages (from pandas>=0.24.0->catboost) (2.8.1)

Requirement already satisfied: retrying>=1.3.3 in c:\users\said\anaconda3\lib\site-packages (from plotly->catboost) (1.3.3)

Requirement already satisfied: cycycler>=0.10 in c:\users\said\anaconda3\lib\site-packages (from matplotlib->catboost) (0.10.0)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\said\anaconda3\lib\site-packages (from matplotlib->catboost) (1.2.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\users\said\anaconda3\lib\site-packages (from matplotlib->catboost) (2.4.7)

```
[2]: #!/usr/bin/env python
```

```
[3]: #
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
import numpy as np
from string import punctuation
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem.snowball import SnowballStemmer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import f1_score
```

```
[4]: data = pd.read_csv('toxic_comments.csv', error_bad_lines=False, engine="python")
data.head()
```

```
[4]:
```

	text	toxic
0	Explanation\nWhy the edits made under my usern...	0
1	D'aww! He matches this background colour I'm s...	0
2	Hey man, I'm really not trying to edit war. It...	0
3	"\nMore\nI can't make any real suggestions on ...	0
4	You, sir, are my hero. Any chance you remember...	0

```
[5]: try:
      data = pd.read_csv('toxic_comments.csv', error_bad_lines=False,
        ↪engine="python")

      except:
          data = pd.read_csv('/datasets/toxic_comments.csv', error_bad_lines=False,
            ↪engine="python")
      data.head()
```

```
[5]:
```

	text	toxic
0	Explanation\nWhy the edits made under my usern...	0
1	D'aww! He matches this background colour I'm s...	0
2	Hey man, I'm really not trying to edit war. It...	0
3	"\nMore\nI can't make any real suggestions on ...	0
4	You, sir, are my hero. Any chance you remember...	0

```
[6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159571 entries, 0 to 159570
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    159571 non-null     object
1    toxic    159571 non-null     int64
dtypes: int64(1), object(1)
memory usage: 2.4+ MB
```

```
[7]: data.columns
```

```
[7]: Index(['text', 'toxic'], dtype='object')
```

```
[8]: import nltk
      nltk.download('punkt')
      nltk.download('stopwords')
```

```
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\said\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\said\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

```
[8]: True
```

```
[9]: import string
      punctuation = string.punctuation
```

```
[10]: noise = stopwords.words('english') + list(punctuation) + list('1234567890')
def tokenizer(value):
    noise = stopwords.words('english') + list(punctuation) + list('1234567890')
    value = value.lower()
    a = word_tokenize(value)
    b = list()
    for el in a:
        if el not in noise:
            if el.isdigit() == False:
                if not el[0].isdigit():
                    b.append(el)
    stemmer = SnowballStemmer('english')
    stemmed_example = [stemmer.stem(w) for w in b]
    a = ' '.join(stemmed_example)
    return a
```

```
[11]: %%time
data['text'] = data['text'].apply(tokenizer)
data.head()
```

Wall time: 9min 35s

```
[11]:
```

	text	toxic
0	explan edit made usernam hardcor metallica fan...	0
1	d'aww match background colour 'm seem stuck th...	0
2	hey man 'm realli tri edit war 's guy constant...	0
3	` ca n't make real suggest improv wonder sect...	0
4	sir hero chanc rememb page 's	0

P.S. , 10

1.2

1.2.1 CountVectorizer Edition

```
[12]: data_cv = data.copy()
x_train, x_val, y_train, y_val = train_test_split(data_cv.drop('toxic', axis=1),
                                                    data_cv['toxic'], test_size=0.
                                                    ↪35, random_state=23)
```

, N 1 2.

```
[13]: count_vec = CountVectorizer(ngram_range=(1,2))
x_train_count = count_vec.fit_transform(x_train['text'])
x_train_count
```

```
[13]: <103721x1760255 sparse matrix of type '<class 'numpy.int64'>'
      with 6061886 stored elements in Compressed Sparse Row format>
```

```
[14]: print(x_train_count.shape)
      print(x_train.shape)
```

```
(103721, 1760255)
(103721, 1)
```

```
[15]: x_val_count = count_vec.transform(x_val['text'])
      x_val_count.shape
```

```
[15]: (55850, 1760255)
```

SVM SGD.

```
[16]: model_sgd1 = SGDClassifier(class_weight='balanced', random_state=131,
      ↪loss='hinge')
      model_sgd1.fit(x_train_count, y_train)
      pred1 = model_sgd1.predict(x_val_count)
      f1_score(y_val, pred1)
```

```
[16]: 0.7818863879957128
```

```
[17]: model_sgd2 = SGDClassifier(class_weight='balanced', random_state=131,
      ↪loss='log')
      model_sgd2.fit(x_train_count, y_train)
      pred2 = model_sgd2.predict(x_val_count)
      f1_score(y_val, pred2)
```

```
[17]: 0.7741500042183415
```

, SVM

SVM SGD

```
[18]: x_train, x_test, y_train, y_test = train_test_split(data_cv.drop('toxic',
      ↪axis=1),
                                                    data_cv['toxic'], test_size=0.
      ↪3, random_state=25433)
```

```
[20]: x_train_count = count_vec.fit_transform(x_train['text'])
      x_test_count = count_vec.transform(x_test['text'])
```

```
[21]: #model_sgd1 = SGDClassifier(class_weight='balanced', random_state=131,
      ↪loss='hinge', n_jobs=2)
```

```
#model_sgd1.fit(x_train, y_train)
model_sgd1.fit(x_train_count, y_train)
predt = model_sgd1.predict(x_test_count)
print('F1 Score on test data:',f1_score(y_test, predt))
```

F1 Score on test data: 0.7887890005288207

1.3

, SGD Classifier hinge loss. Countvectorizer N
(1,1),
SVM , . ,

1.4 -

- ☒ Jupyter Notebook
- ☒
- ☒
- ☒
- ☒
- ☒ $F1$ 0.75
- ☒

[]: